



OpenDS ZKI AK Verzeichnisdienste Oct. 2009

Abdi Mohammadi

Principal Field Technologist
Sun Microsystems, Inc.



Agenda

- Introduction to OpenDS
- Features
- Replication Topology
- Embedded OpenDS
- Performance
 - > Scaling from Entreprises to Service Providers
- MySQL NDB as Database Backend
- Roadmap & Integration in DSEE

Agenda

- **Introduction to OpenDS**
- Features
- Replication Topology
- Embedded OpenDS
- Performance
 - > Scaling from Entreprises to Service Providers
- MySQL NDB as Database Backend
- Roadmap & Integration in DSEE

The OpenDS project



- Released in Open Source in July 2006
 - > CDDL
 - > Source code at <https://opens.dev.java.net/>
- Sponsored by Sun Microsystems
- A Java based Server supporting the LDAPv3 protocol based on Berkeley DB Java Edition
 - > Not accessible from outside
- All security, access controls, password management features to safely store the information about Users

What is it for ?

- Generic object oriented data store
- White pages and Email Address Book
- Mostly the data store for Identities
 - > For Authentication and Authorization
 - > For profiles and personalization
- The underlying infrastructure in all Enterprises
 - > Leveraged by Web and Mail infrastructure products
 - > Cornerstone of Identity Management products:
 - > Access Management and Federation
 - > Provisioning and De-provisioning tools

OpenDS Goals

- A complete set of Directory Services
 - > Directory Back-end database
 - > Full LDAPv3 compliance and standard extensions
 - > Multi-Master replication
 - > Directory Proxy Services : load-balancing, data distribution, security services
 - > Virtual Directory Capabilities
- Horizontal and Vertical Scalability
- Sun Directory Server Enterprise Edition 8.0 will be based on OpenDS code

Three Principles

- Ease of Use
 - > Installation, Configuration, Management, Monitoring...
- Performance
 - > Code with performance in mind:
 - > Careful with memory
 - > Threads and reduced contention
 - > Monitors and Configurable queues
- Extensibility
 - > Many interfaces defined
 - > Default implementation provided

Agenda

- Introduction to OpenDS
- **Features**
- Replication Topology
- Embedded OpenDS
- Performance
 - > Scaling from Entreprises to Service Providers
- MySQL NDB as Database Backend
- Roadmap & Integration in DSEE

Features

- LDAPv3 directory server fully standard compliant
 - > Supports many LDAP standard and experimental extensions
 - > Supports Multi Master Replication with 3 different levels of data consistency
 - > Extensive security features
 - > StartTLS, SASL, Password Policy, ACIs
- Virtual Attributes such as isMemberOf
- APIs to add components, plugins, controls, extensions and extended operations to the server

Features (cont.)

- DSML Gateway
- Graphical Control-Panel and Command Line Tools for configuration or monitoring
- Complete User Administration and Reference documentation
- Localized in 6 different languages
 - > Community lead translation, You contribute
- Support available with Sun OpenDS Standard Edition 2.0.
- Installs in 6 clicks and less than 3 minutes

<https://www.opensds.org/promoted-builds/latest/install/QuickSetup.jnlp>

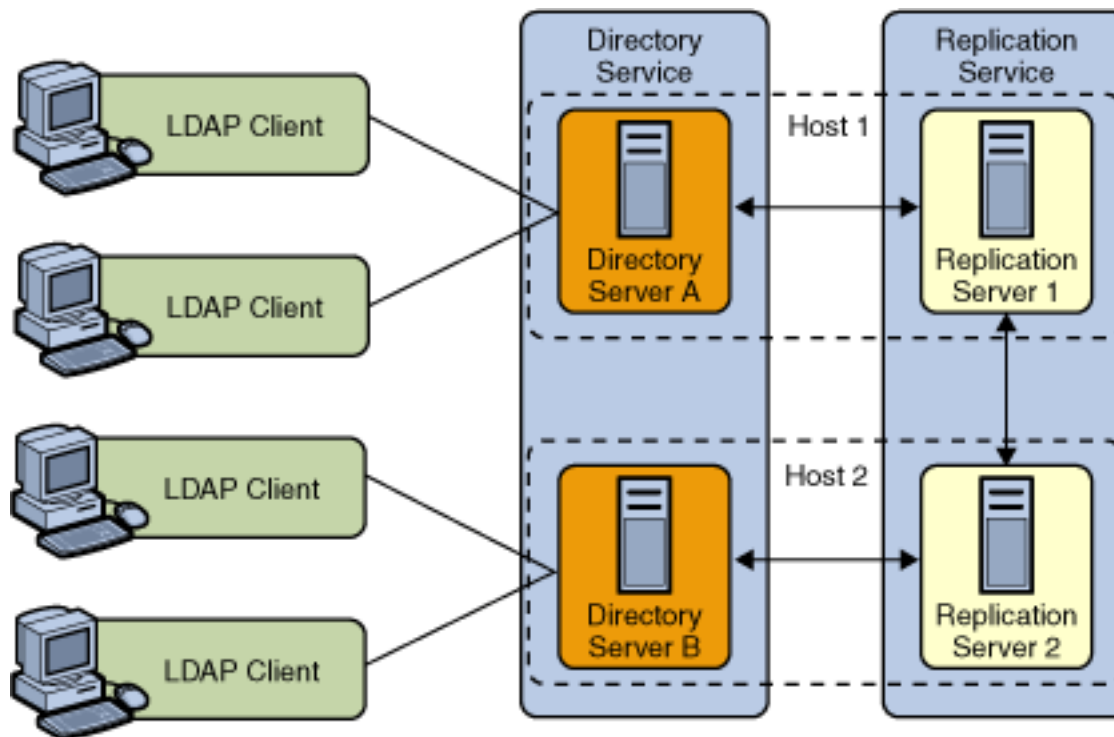
Agenda

- Introduction to OpenDS
- Features
- **Replication Topology**
- Embedded OpenDS
- Performance
 - > Scaling from Entreprises to Service Providers
- MySQL NDB as Database Backend
- Roadmap & Integration in DSEE

Replication Features

- N-Way Multi Master
- Schema Replication
- Fractional Replication
- Strong Authentication
- WAN Support

Basic Replication Architecture



Replication Servers

A replication server performs the following tasks:

- Manages connections from directory servers
- Connects to other replication servers
- Listens for connections from other replication servers
- Receives changes from directory servers
- Forwards changes to directory servers and other replication servers
- Saves changes to stable storage, which includes trimming older operations


Centralized replication service

- Simpler topology
- Too many replication agreements hurt performances.
- Easier to implement assured, repair monitoring
- Easier to deploy large topologies
- All servers are writable.
- WAN Support

Replication Groups

Designed to support multi-data center deployments and disaster recovery scenarios

Group IDs determine how a directory server domain connects to an available replication server



Disaster Recovery deployment: 2 data centers with different group ids

Assured Replication

- Asynchronous Replication but wait for data to be replicated before sending ACK to application.
- Two modes :
 - > Safe Data :
Save data on Replication Servers before returning ACK to the LDAP application.
 - > Safe Read :
Propagate Change to selected Directory Servers before returning ACK to the LDAP application.
- <https://opends.org/wiki/page/AssuredReplication>

Agenda

- Introduction to OpenDS
- Features
- Replication Topology
- **Embedded OpenDS**
- Performance
 - > Scaling from Entreprises to Service Providers
- MySQL NDB as Database Backend
- Roadmap & Integration in DSEE

Embedding OpenDS in Apps

- OpenDS server runs in the same JVM as your application, standalone or web application
- OpenDS jars and files are becoming part of your application
- The applications controls when OpenDS starts and stops
- Application can access the server either
 - > Using the LDAP protocol over the network
 - > Or using Internal APIs
- Still can use Replication between instances

Why embedding ?

- To provide a better “Out of the Box” experience and secure access to the data
- To reduce administration and initial configuration
- To setup an optimized service for your app's needs
- To reduce memory footprint by running in the same JVM
- To deploy in Web Application Container
- To test simple automatic testing of LDAP client code

Cost and Limitations

- OpenDS 2.0 Jars are less than 24 MB
- OpenDS configuration and schema is ~ 5MB
- Need to tune JVM Heap size accordingly for both application and OpenDS services
- Disable specific OpenDS services if unused
- Requires local disk and performances are mostly tied to the storage system
- Scalability: limited by the amount of memory available.

Example use



- OpenSSO
 - > Embeds OpenDS as its configuration store
 - > Replicates the configuration between instances for HA
 - > Can be used as the User store (though not officially supported)
 - > In a short future, will also store XACML policies for the Policy engine.
 - > Goal is to scale up to 1 000 000 policies
 - > Using Internal APIs for better performances

<https://www.openss.org/wiki/page/UsingOpenDSAsAnEmbeddedApplication>

https://www.openss.org/wiki/attach/OpenDSPresentations/OpenDS_Jazoon08.pdf

Agenda

- Introduction to OpenDS
- Features
- Replication Topology
- Embedded OpenDS
- **Performance**
 - > **Scaling from Enterprise to Service Providers**
- MySQL NDB as Database Backend
- Roadmap & Integration in DSEE

Scalability

- A single server scales from few entries to several tens of millions of entries.
- No hard limit
 - > Limit is based on performance requirements and Service Level agreements for specific hardware (i.e. cost)
 - > How long to import
 - > Response time, Throughput...
- Scale Read operations by Replicating the server
- Scale Write operations by distributing the data onto multiple servers

Performances

- LDAP servers performances are relative to:
 - > The data
 - > Access patterns (connections, reads, writes, ...)
 - > CPUs, memory, filesystems and storage subsystem
 - > Many more parameters...
- OpenDS baseline:
 - > 10 Millions entries database (avg 1.5K/entry)
 - > Can achieve up to 100K Search/Sec, 12K MODs /Sec in real deployment scenario
 - > Response time is below 1ms for both reads and writes.

Performance (cont.)

- We managed to get more than 100 000 searches / seconds on a single box with 10 million entries
 - > Intel Core i7 processors (8 core)
 - > Machine over-clocked
 - > Real numbers at :
http://blogs.sun.com/ds/entry/opens_nehalem_benchmark_sunblade_6270
- Throughput is only one measure:
 - > What is the average response time ?
 - > What is the maximum response time ?

Performance (cont.)

- **The Context**

- > Sun X4150
- > 8 x Intel 3.2GHz
- > 64GB RAM

- **Search rate**

- > 8 clients / CPU 35% idle
- > 15500 op/s
- > 10% = 0.193417
- > 50% = 0.223053
- > 90% = 0.278756
- > 99% = 0.362329
- > 99.9% = 0.422575
- > 99.99% = 35.5056
- > 99.999% = 41.8817
- > **Average = 0.237412 ms**

- > Internal disk
- > 10M 1.5K entries
- > Fully preloaded

- **Modify rate**

- > 2 clients / CPU 75% idle
- > 4000 op/s
- > 10% = 0.237901
- > 50% = 0.288164
- > 90% = 0.36565
- > 99% = 0.486679
- > 99.9% = 0.706433
- > 99.99% = 11.1529
- > 99.999% = 65.5304
- > **Average = 0.303045 ms**

Who Needs this ?

- Telecom operators have used Directory services for Business side for years (Portal...)
 - > Storing identities of their Customers and phones
 - > Largest known deployment (with Sun Directory Server) is 120 M entries
- OpenDS for Naming Services
 - > Solaris and Linux
 - <http://developers.sun.com/identity/reference/techart/opens-namesvcs.html>
 - > SAMBA <https://www.opens.org/wiki/page/SambaCIFSServer>
 - > Kerberos <https://www.opens.org/wiki/page/ConfiguringSASLGSSAPIAuthentication>

Agenda

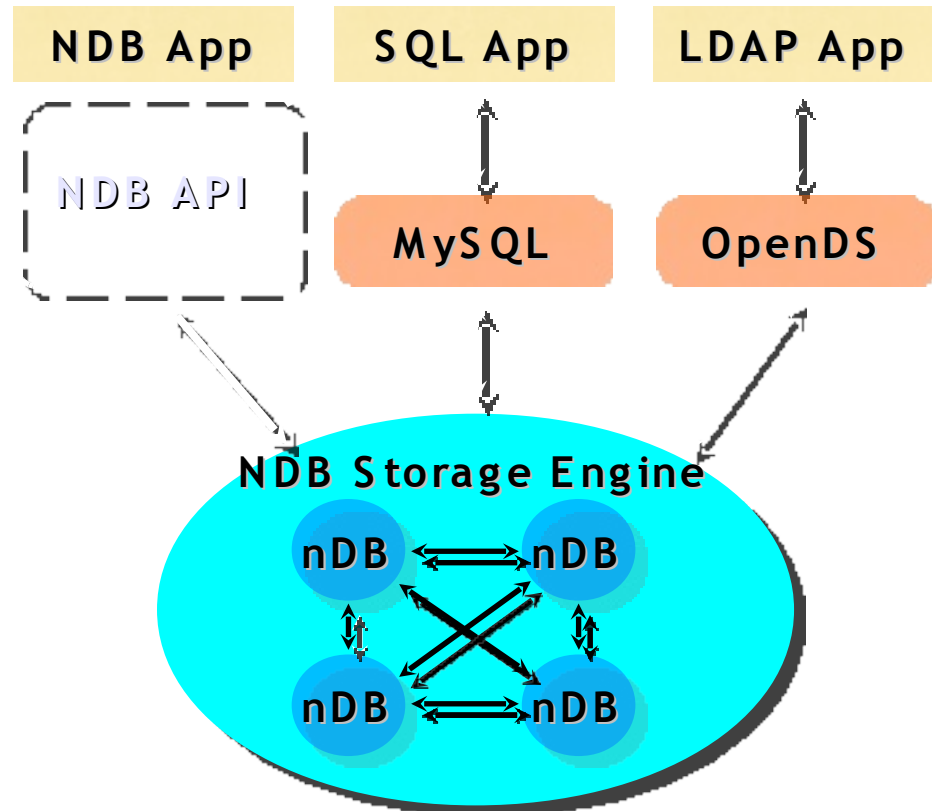
- Introduction to OpenDS
- Features
- Replication Topology
- Embedded OpenDS
- Performance
 - > Scaling from Entreprises to Service Providers
- **MySQL NDB as Database Backend**
- Roadmap & Integration in DSEE

MySQL Cluster

- It's the Carrier Grade edition of MySQL database product
 - > Provides 99,999% availability
 - > Runs on commodity hardware
 - > Performances are constant
- Uses Network DB (NDB):
 - > Fully transactional
 - > Partition the data across multiple nodes
 - > Memory Based for Performances
 - > Provides automatic replication and fail-over

OpenDS and MySQL Cluster

- Simultaneous access to the Data via LDAP, SQL or direct API
- Direct access to NDB
- Need LDAP Meta Data to be in NDB
- Implemented in OpenDS as an optional back-end



OpenDS NDB Back-end

- Uses NDB/J bindings to NDB API (JNI)
- Allows multiple OpenDS process to access the same NDB database concurrently
- Maps the SQL Tables schema to LDAP Schema
- Requires additional tables in NBD
 - > For DN's, Objectclasses and other LDAP specific data

OpenDS NDB Back-end (cont.)

- OpenDS/NDB is feature-complete
- Finalizing perf and stress tests
- Code is available in OpenDS Trunk
- How To :
<https://www.opens.org/wiki/page/EnableNDBBackend>



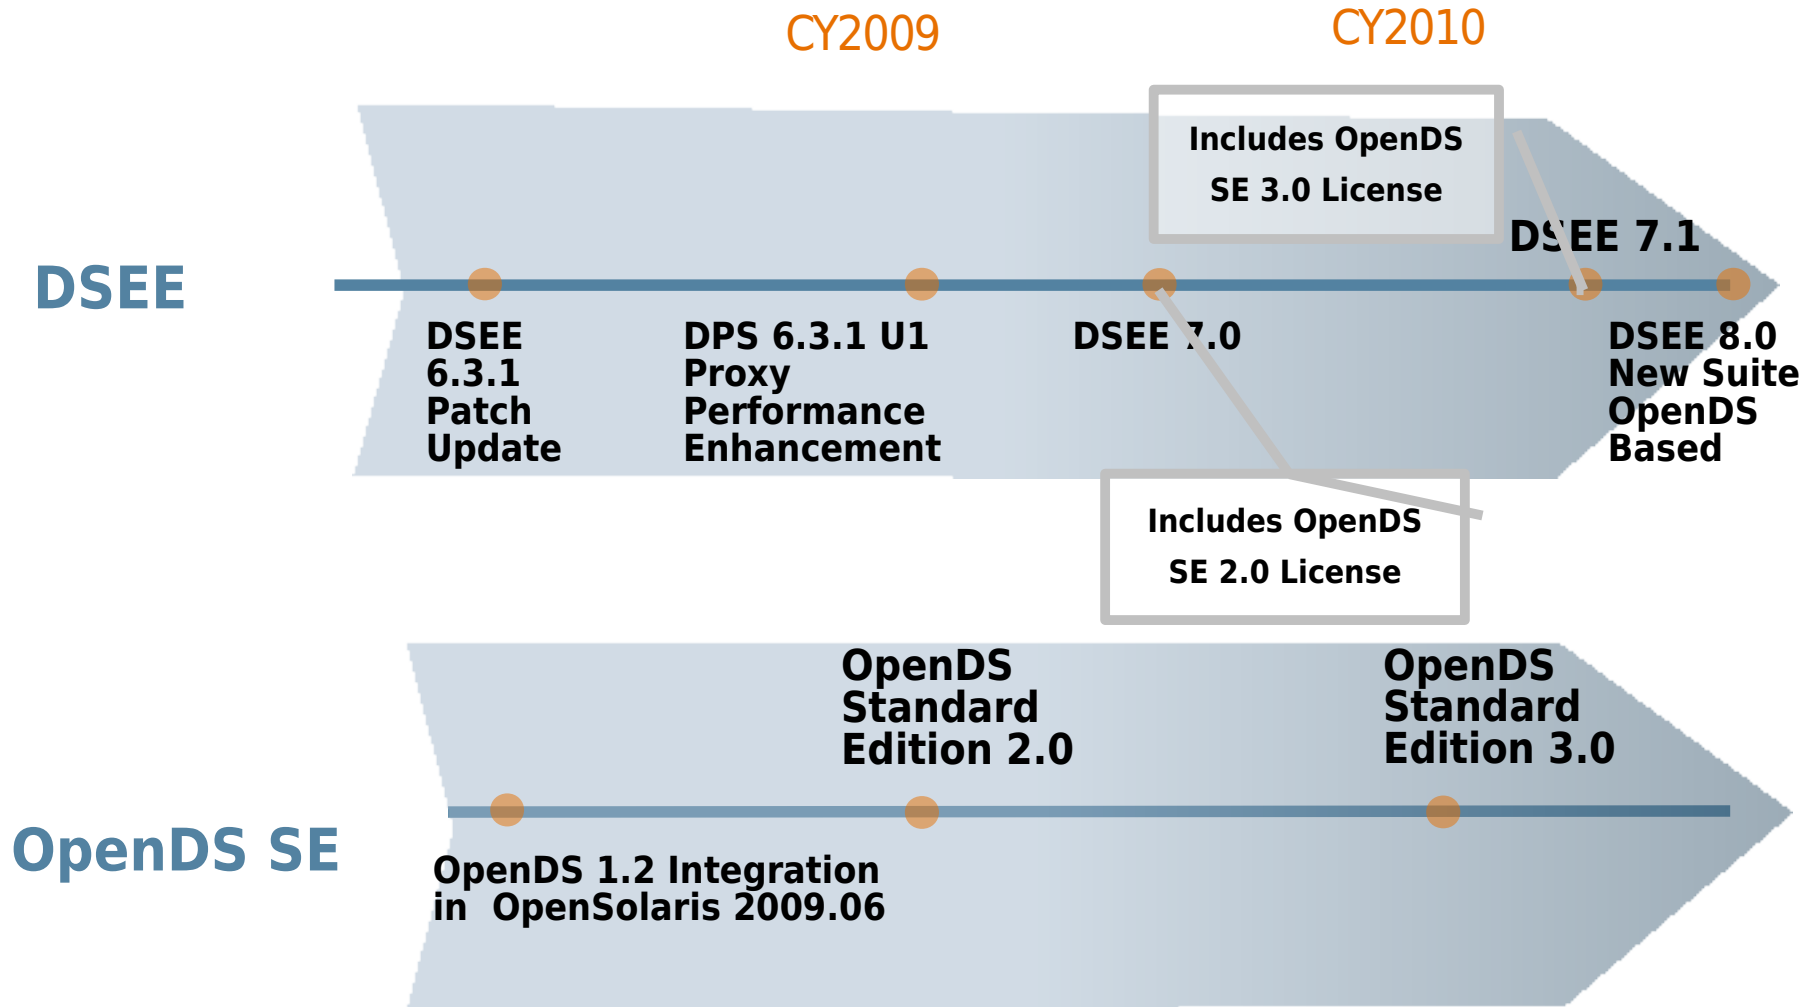
Benefits

- Lower cost
 - > Runs on commodity hardware
 - > Single database for multiple LDAP and SQL servers
- Performance scales linearly
 - > Add more database nodes to scale capacity
 - > Add more OpenDS server to scale throughput
- Response time is higher than a single node, but that's not the point

Agenda

- Introduction to OpenDS
- Features
- Replication Topology
- Embedded OpenDS
- Performance
 - > Scaling from Entreprises to Service Providers
- MySQL NDB as Database Backend
- **Roadmap & Integration in DSEE**

Directory Services Roadmap



* = All future dates subject to change, all calendar year

Next Steps...

- Give OpenDS a try:
 - > <http://www.opens.org>
- Participate in the community
 - > Join/Login on Java.net
 - > Request a Role in the OpenDS project
 - > <http://opens.dev.java.net>
 - > IRC: #opens on freenode.net
- Help with Translation to your preferred language





Thank You !

**OpenDS
ZKI AK Verzeichnisdienste
Oct. 2009**

Abdi.Mohammadi@sun.com